Assignment 1 - Introduction to Data Science and Python

November 1, 2024

Stefan Dimitrov Velev, 0MI3400521, Big Data Technologies

Faculty of Mathematics and Informatics, Sofia University

1. Import required Python packages

[1]:
```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

2. Read the CSV file - Life expectancy vs. GDP per capita, 2022

Data source: UN, World Population Prospects (2024); World Bank (2023)

https://ourworldindata.org/grapher/life-expectancy-un-vs-gdp-per-capita-wb

[2]:
```
df = pd.read_csv('./data/life-expectancy-un-vs-gdp-per-capita-wb.csv',␣
↪delimiter=',')
```

[3]:
```
df.head()
```

[3]:
```
        Entity      Code   Year  \
0     Abkhazia  OWID_ABK   2015
1  Afghanistan       AFG -10000
2  Afghanistan       AFG  -9000
3  Afghanistan       AFG  -8000
4  Afghanistan       AFG  -7000

   Life expectancy - Sex: all - Age: 0 - Variant: estimates  \
0                                                NaN
1                                                NaN
2                                                NaN
3                                                NaN
4                                                NaN

   GDP per capita, PPP (constant 2017 international $)  \
0                                                NaN
1                                                NaN
2                                                NaN
3                                                NaN
4                                                NaN
```

```
   Population (historical) Continent
0                      NaN      Asia
1                  14737.0       NaN
2                  20405.0       NaN
3                  28253.0       NaN
4                  39120.0       NaN
```

[4]:
```
print("The number of rows in the data frame is:", len(df))
```

```
The number of rows in the data frame is: 59858
```

3. Data Cleaning

[5]:
```
# Remove the unnecessary columns in the data frame
df = df[['Entity', 'Year', 'Life expectancy - Sex: all - Age: 0 - Variant:␣
 ↪estimates', 'GDP per capita, PPP (constant 2017 international $)',␣
 ↪'Population (historical)']]
```

[6]:
```
# Rename the applicable columns
df = df.rename(columns={'Entity': 'Country', 'Life expectancy - Sex: all - Age:␣
 ↪0 - Variant: estimates': 'Life expectancy', 'GDP per capita, PPP (constant␣
 ↪2017 international $)': 'GDP per capita', 'Population (historical)':␣
 ↪'Population'})
```

[7]:
```
# Leaving only rows for year 2022
df = df[df['Year'] == 2022]
```

[8]:
```
# Remove rows with missing values
df = df.dropna()
```

[9]:
```
# Remove not-country-specific entries
df = df[df['Country'] != 'High-income countries']
df = df[df['Country'] != 'Low-income countries']
df = df[df['Country'] != 'Lower-middle-income countries']
df = df[df['Country'] != 'Upper-middle-income countries']
df = df[df['Country'] != 'World']
```

[10]:
```
df[df['Country'] == 'Central African Republic']
```

[10]:
```
                       Country  Year  Life expectancy  GDP per capita  \
10140  Central African Republic  2022           18.818        823.9822

       Population
10140   5098038.0
```

[11]:
```
# Correct the Central African Republic life expectancy according the World Bank␣
 ↪Report for 2022
# Source: https://data.worldbank.org/indicator/SP.DYN.LE00.IN
```

```
df.loc[df['Country'] == 'Central African Republic', 'Life expectancy'] = 54.48
```

[12]:
```
df.count()
```

[12]:
```
Country            188
Year               188
Life expectancy    188
GDP per capita     188
Population         188
dtype: int64
```

4. Draw a scatter plot of GDP per capita vs Life expectancy for 2022

[13]:
```
df.plot.scatter(x='GDP per capita', y='Life expectancy')
plt.xlabel('GDP per Capita (international-$ in 2017 prices)')
plt.ylabel('Life Expectancy (years)')
plt.title('GDP per Capita vs Life Expectancy (world countries) for 2022')
plt.show()
```

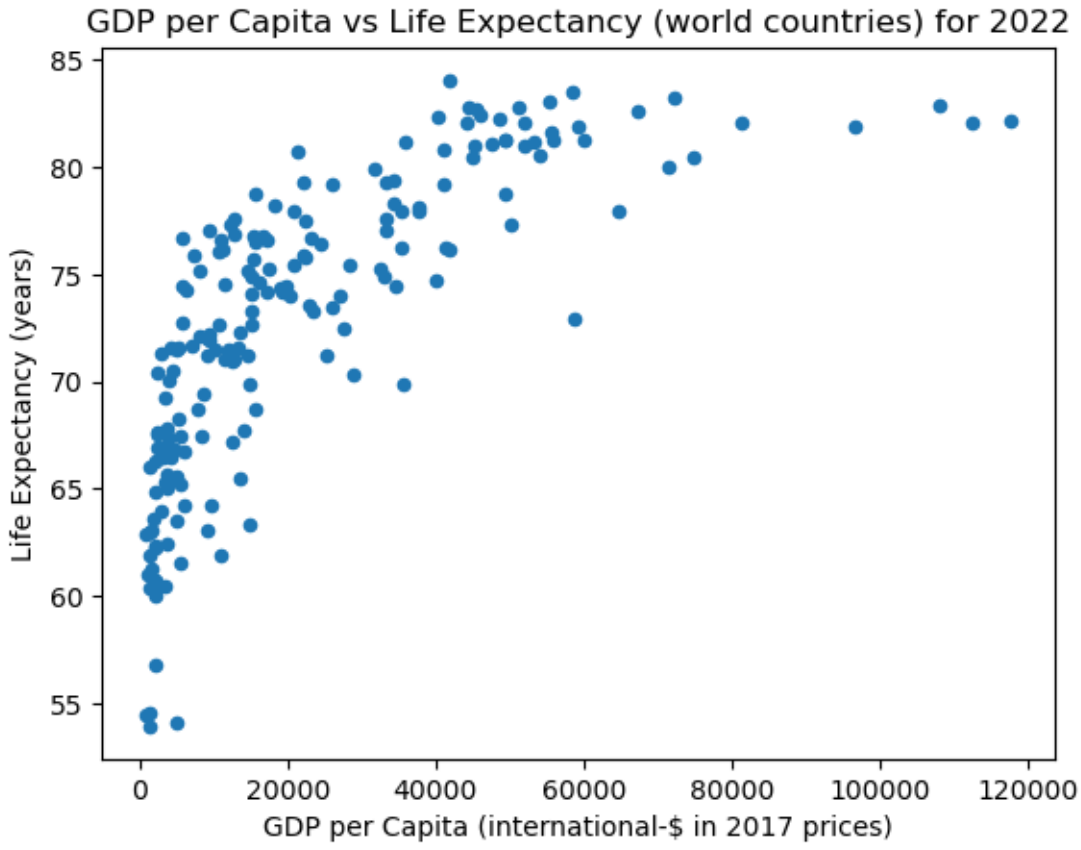

[14]:
```
plt.figure(figsize=(18, 8))
plt.scatter(df['GDP per capita'], df['Life expectancy'], color='blue', s =␣
 ↪df['Population']/1000000, alpha=0.5)

for i, country in enumerate(df['Country']):
    plt.text(df['GDP per capita'].iloc[i], df['Life expectancy'].iloc[i],␣
 ↪df['Country'].iloc[i], fontsize=6, ha='right')

plt.xticks([0, 10000, 20000, 30000, 40000, 50000, 60000, 70000, 80000, 90000,␣
 ↪100000, 110000, 120000])
plt.xlabel('GDP per Capita (international-$ in 2017 prices)')
plt.ylabel('Life Expectancy (years)')
plt.title('GDP per Capita vs Life Expectancy (world countries) for 2022')
plt.show()
```

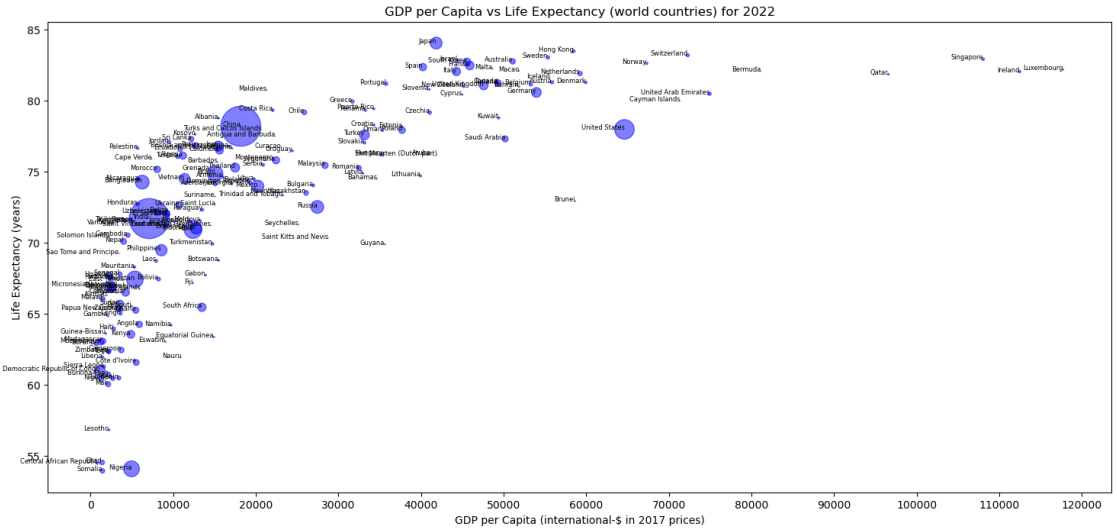

[15]:
```
plt.figure(figsize=(18, 8))
df_lower_GDP = df[df['GDP per capita'] <= 20000]
plt.scatter(df_lower_GDP['GDP per capita'], df_lower_GDP['Life expectancy'],␣
 ↪color='blue', s = df_lower_GDP['Population']/200000, alpha=0.5)

for i, country in enumerate(df_lower_GDP['Country']):
    plt.text(df_lower_GDP['GDP per capita'].iloc[i], df_lower_GDP['Life␣
 ↪expectancy'].iloc[i], df_lower_GDP['Country'].iloc[i], fontsize=6,␣
 ↪ha='right')

plt.ylim(53, 86)

plt.xlabel('GDP per Capita (international-$ in 2017 prices)')
plt.ylabel('Life Expectancy (years)')
```

```
plt.title('GDP per Capita vs Life Expectancy (countries with GDP less than 20␣
 ↪000) for 2022')
plt.show()
```

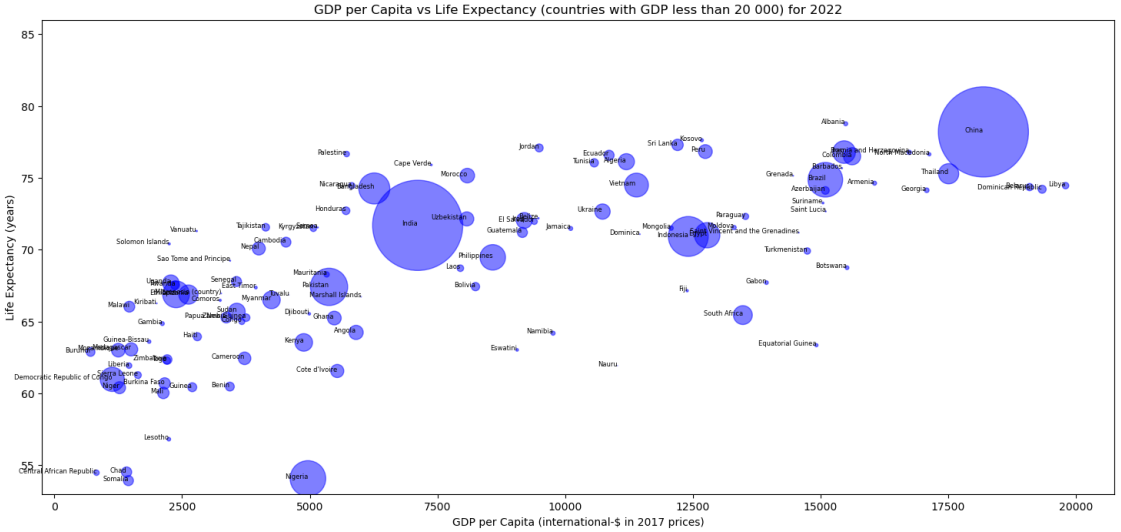

[16]:
```
plt.figure(figsize=(18, 8))
df_higher_GDP = df[df['GDP per capita'] >= 20000]
plt.scatter(df_higher_GDP['GDP per capita'], df_higher_GDP['Life expectancy'],␣
 ↪color='blue', s = df_higher_GDP['Population']/200000, alpha=0.5)

for i, country in enumerate(df_higher_GDP['Country']):
    plt.text(df_higher_GDP['GDP per capita'].iloc[i], df_higher_GDP['Life␣
 ↪expectancy'].iloc[i], df_higher_GDP['Country'].iloc[i], fontsize=6,␣
 ↪ha='right')

plt.ylim(53, 86)

plt.xlabel('GDP per Capita (international-$ in 2017 prices)')
plt.ylabel('Life Expectancy (years)')
plt.title('GDP per Capita vs Life Expectancy (countries with GDP more than 20␣
 ↪000) for 2022')
plt.show()
```

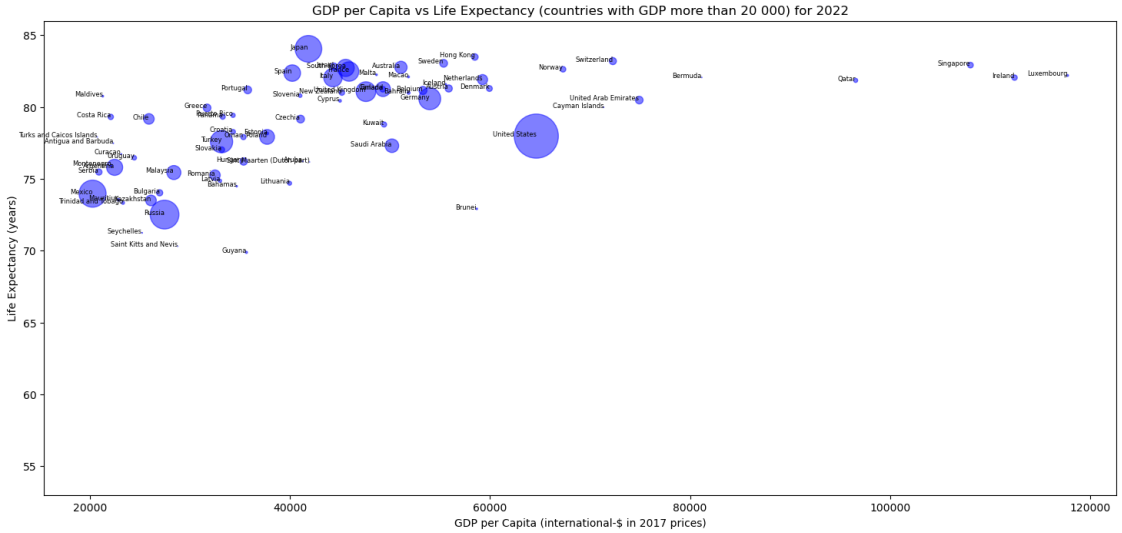

[17]:
```
plt.figure(figsize=(18, 8))
plt.scatter(df_higher_GDP['GDP per capita'], df_higher_GDP['Life expectancy'],␣
 ↪color='blue', s = df_higher_GDP['Population']/200000, alpha=0.5)

for i, country in enumerate(df_higher_GDP['Country']):
    plt.text(df_higher_GDP['GDP per capita'].iloc[i], df_higher_GDP['Life␣
 ↪expectancy'].iloc[i], df_higher_GDP['Country'].iloc[i], fontsize=6,␣
 ↪ha='right')

plt.xlabel('GDP per Capita (international-$ in 2017 prices)')
plt.ylabel('Life Expectancy (years)')
plt.title('GDP per Capita vs Life Expectancy (countries with GDP more than 20␣
 ↪000) for 2022')
plt.show()
```

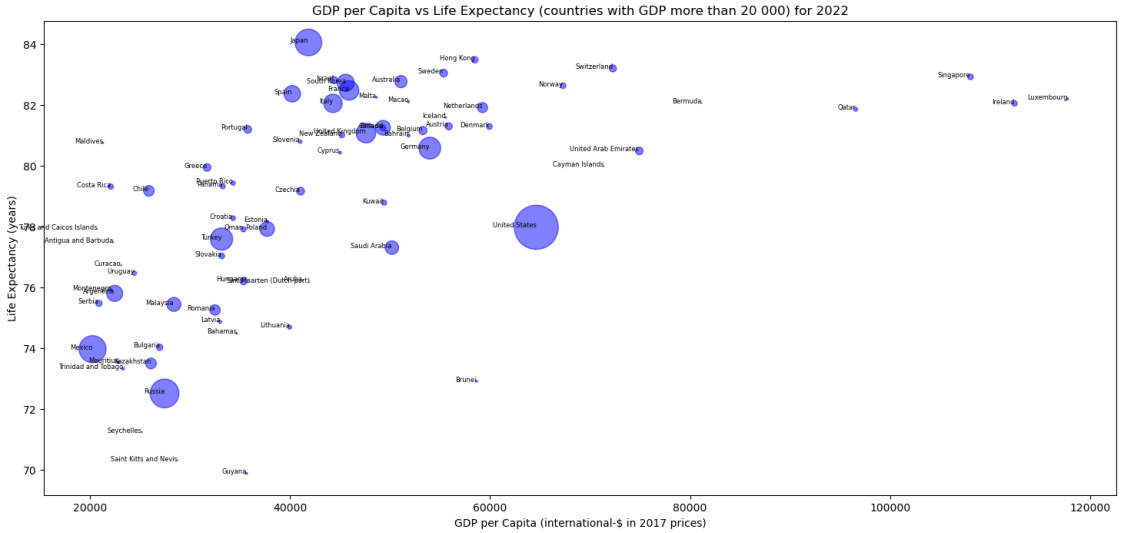

5. Find out which countries have a life expectancy higher than one standard deviation above the mean?

[18]:
```
# Get statistic for the whole DataFrame with Pandas
df.describe()
```

[18]:

	Year	Life expectancy	GDP per capita	Population
count	188.0	188.000000	188.000000	1.880000e+02
mean	2022.0	72.752750	22643.151954	4.151910e+07
std	0.0	7.039319	22790.551663	1.522105e+08
min	2022.0	53.931000	708.178300	1.001200e+04
25%	2022.0	67.404750	5134.252600	1.831064e+06
50%	2022.0	74.079500	14829.153500	7.917908e+06
75%	2022.0	77.937000	34831.845750	3.040557e+07
max	2022.0	84.054000	117746.990000	1.425423e+09

[19]:
```
# Get the average life expectancy and the standard deviation with NumPy
mean_life_expectancy = np.mean(df['Life expectancy'])
standard_deviation_life_expectancy = np.std(df['Life expectancy'])
print('Mean life expectancy is', mean_life_expectancy)
print('Standard deviation of life expectancy is',␣
 ↪standard_deviation_life_expectancy)
```

```
Mean life expectancy is 72.75275
Standard deviation of life expectancy is 7.020572922933241
```

[20]:
```
# Find lower boundary for the searched countries - one standard deviation above␣
 ↪the mean
lower_boundary_high_life_expectancy = mean_life_expectancy +␣
 ↪standard_deviation_life_expectancy
```

```
print('The searched lower boundary for high life expectancy is',␣
 ↪lower_boundary_high_life_expectancy)
```

```
The searched lower boundary for high life expectancy is 79.77332292293325
```

[21]:
```
# List countries with higher life expectancy (one standard deviation above the␣
 ↪mean)
df_higher_life_expectancy = df[df['Life expectancy'] >␣
 ↪lower_boundary_high_life_expectancy]
print('The number of countries with higher life expectancy is',␣
 ↪len(df_higher_life_expectancy))
print('The list of countries with higher life expectancy is',  ', '.
 ↪join(df_higher_life_expectancy['Country']))
```

```
The number of countries with higher life expectancy is 36
The list of countries with higher life expectancy is Australia, Austria,
Bahrain, Belgium, Bermuda, Canada, Cayman Islands, Cyprus, Denmark, Finland,
France, Germany, Greece, Hong Kong, Iceland, Ireland, Israel, Italy, Japan,
Luxembourg, Macao, Maldives, Malta, Netherlands, New Zealand, Norway, Portugal,
Qatar, Singapore, Slovenia, South Korea, Spain, Sweden, Switzerland, United Arab
Emirates, United Kingdom
```

6. Check which countries have high life expectancy but have low GDP?

[22]:
```
# List countries with high life expectancy but low GDP using means
mean_gdp = np.mean(df['GDP per capita'])
print('Mean GDP per capita is', mean_gdp)
df_higher_life_expectancy_lower_GDP1 = df[(df['Life expectancy'] >␣
 ↪mean_life_expectancy) & (df['GDP per capita'] < mean_gdp)]
print('The number of countries with high life expectancy but low GDP per capita␣
 ↪using means is', len(df_higher_life_expectancy_lower_GDP1))
print('The list of countries with high life expectancy but low GDP per capita␣
 ↪using means is:',  ', '.
 ↪join(df_higher_life_expectancy_lower_GDP1['Country']))
```

```
Mean GDP per capita is 22643.151954255318
The number of countries with high life expectancy but low GDP per capita using
means is 38
The list of countries with high life expectancy but low GDP per capita using
means is: Albania, Algeria, Antigua and Barbuda, Argentina, Armenia, Azerbaijan,
Bangladesh, Barbados, Belarus, Bosnia and Herzegovina, Brazil, Cape Verde,
China, Colombia, Costa Rica, Dominican Republic, Ecuador, Georgia, Grenada,
Iran, Jordan, Kosovo, Libya, Maldives, Mexico, Montenegro, Morocco, Nicaragua,
North Macedonia, Palestine, Peru, Serbia, Sri Lanka, Suriname, Thailand,
Tunisia, Turks and Caicos Islands, Vietnam
```

[23]:
```
plt.figure(figsize=(18, 8))
```

```
plt.scatter(df_higher_life_expectancy_lower_GDP1['GDP per capita'],␣
 ↪df_higher_life_expectancy_lower_GDP1['Life expectancy'], color='blue', s =␣
 ↪df_higher_life_expectancy_lower_GDP1['Population']/200000, alpha=0.5)

for i, country in enumerate(df_higher_life_expectancy_lower_GDP1['Country']):
    plt.text(df_higher_life_expectancy_lower_GDP1['GDP per capita'].iloc[i],␣
 ↪df_higher_life_expectancy_lower_GDP1['Life expectancy'].iloc[i],␣
 ↪df_higher_life_expectancy_lower_GDP1['Country'].iloc[i], fontsize=6,␣
 ↪ha='right')

plt.xlabel('GDP per Capita (international-$ in 2017 prices)')
plt.ylabel('Life Expectancy (years)')
plt.title('GDP per Capita vs Life Expectancy (countries with high life␣
 ↪expectancy but low GDP per Capita using means as a criteria) for 2022')
plt.show()
```

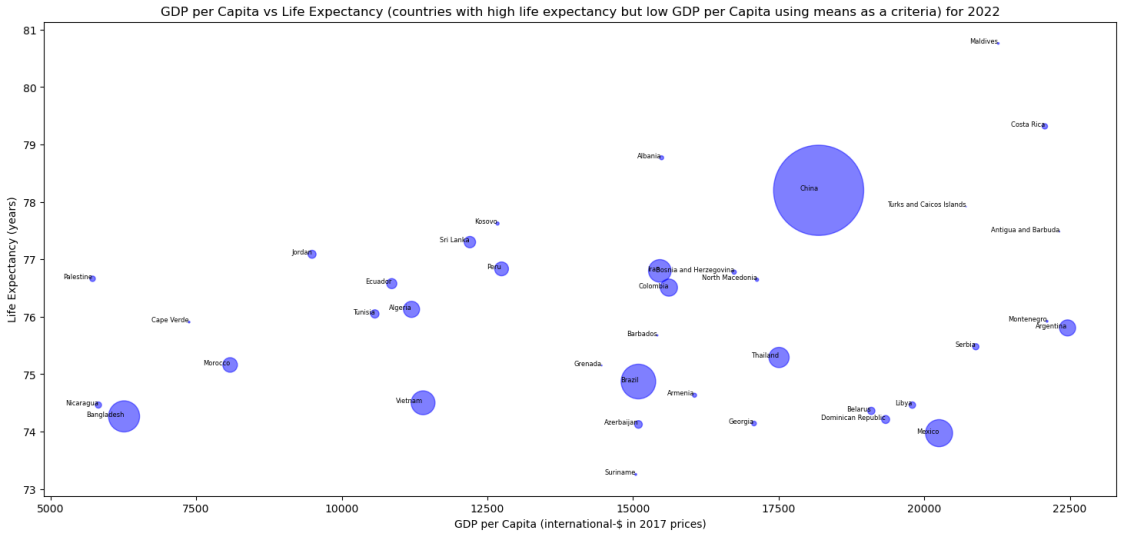

[24]:
```
# List countries with high life expectancy but low GDP using 75th percentile␣
 ↪for high life expectancy and the 25th percentile for low GDP
upper_boundary_low_gdp = df['GDP per capita'].quantile(0.25)
print('The upper boundary for low GDP per capita is', upper_boundary_low_gdp)
lower_boundary_high_life_expectancy = df['Life expectancy'].quantile(0.75)
print('The lower boundary for high life expectancy is',␣
 ↪lower_boundary_high_life_expectancy)
df_higher_life_expectancy_lower_GDP2 = df[(df['Life expectancy'] >=␣
 ↪lower_boundary_high_life_expectancy) & (df['GDP per capita'] <=␣
 ↪upper_boundary_low_gdp)]
print('The number of countries with high life expectancy but low GDP per capita␣
 ↪using percentiles is', len(df_higher_life_expectancy_lower_GDP2))
```

```
The upper boundary for low GDP per capita is 5134.2526
The lower boundary for high life expectancy is 77.937
The number of countries with high life expectancy but low GDP per capita using
percentiles is 0
```

7. Find whether each strong economy (normally indicated by GDP) have high life expectancy?

[25]:
```
lower_boundary_high_gdp = df['GDP per capita'].quantile(0.75)
print('The lower boundary for high GDP per capita is', lower_boundary_high_gdp)
df_higher_gdp = df[df['GDP per capita'] >= lower_boundary_high_gdp]
print('The number of countries with high GDP per capita using 75th percentile␣
 ↪is', len(df_higher_gdp))
```

```
The lower boundary for high GDP per capita is 34831.84575
The number of countries with high GDP per capita using 75th percentile is 47
```

[26]:
```
df_higher_gdp
```

[26]:

	Country	Year	Life expectancy	GDP per capita	\
2853	Aruba	2022	76.226	41273.613	
3579	Australia	2022	82.766	51090.260	
3840	Austria	2022	81.296	55867.184	
4624	Bahrain	2022	80.992	51854.715	
5661	Belgium	2022	81.159	53287.152	
6273	Bermuda	2022	82.062	81165.650	
7998	Brunei	2022	72.917	58669.902	
9564	Canada	2022	81.249	49296.380	
9879	Cayman Islands	2022	79.984	71353.890	
13156	Cyprus	2022	80.434	44996.316	
13417	Czechia	2022	79.165	41052.348	
14024	Denmark	2022	81.291	59935.120	
16671	Estonia	2022	78.167	37711.820	
18745	Finland	2022	81.243	49275.152	
19006	France	2022	82.475	45904.410	
20258	Germany	2022	80.580	53969.625	
22553	Guyana	2022	69.888	35634.688	
23603	Hong Kong	2022	83.485	58478.883	
23864	Hungary	2022	76.212	35356.777	
24125	Iceland	2022	81.588	55567.438	
25430	Ireland	2022	82.050	112445.420	
25767	Israel	2022	82.814	44393.300	
26028	Italy	2022	82.052	44292.190	
26542	Japan	2022	84.054	41837.910	
27963	Kuwait	2022	78.788	49400.355	
30715	Lithuania	2022	74.696	39955.246	
31506	Luxembourg	2022	82.201	117746.990	
31592	Macao	2022	82.103	51840.140	
33118	Malta	2022	82.250	48641.850	

37552	Netherlands	2022	81.912	59249.168
38016	New Zealand	2022	81.006	45185.312
40106	Norway	2022	82.631	67296.160
40706	Oman	2022	77.911	35336.895
42983	Poland	2022	77.923	37706.605
43244	Portugal	2022	81.194	35767.723
43624	Qatar	2022	81.857	96557.810
46611	Saudi Arabia	2022	77.310	50188.297
47945	Singapore	2022	82.921	108036.110
48019	Sint Maarten (Dutch part)	2022	76.180	41942.918
48541	Slovenia	2022	80.793	41015.227
50005	South Korea	2022	82.727	45560.125
50493	Spain	2022	82.366	40223.010
51673	Sweden	2022	83.046	55359.344
51934	Switzerland	2022	83.200	72278.210
56201	United Arab Emirates	2022	80.487	74917.670
56462	United Kingdom	2022	81.074	47587.168
56723	United States	2022	77.979	64623.125

	Population
2853	107792.0
3579	26200987.0
3840	9064679.0
4624	1533459.0
5661	11641813.0
6273	64772.0
7998	455374.0
9564	38821260.0
9879	71609.0
13156	1331376.0
13417	10673216.0
14024	5902898.0
16671	1350092.0
18745	5569299.0
19006	66277412.0
20258	84086228.0
22553	821636.0
23603	7465914.0
23864	9684306.0
24125	380368.0
25430	5110013.0
25767	9103144.0
26028	59619106.0
26542	124997586.0
27963	4589514.0
30715	2816922.0
31506	653316.0

```
31592     704359.0
33118     528194.0
37552   17904422.0
38016    5131733.0
40106    5456795.0
40706    4730227.0
42983   38385734.0
43244   10417075.0
43624    2892465.0
46611   32175352.0
47945    5649886.0
48019      42163.0
48541    2115230.0
50005   51782514.0
50493   47828386.0
51673   10487333.0
51934    8792180.0
56201   10242085.0
56462   68179315.0
56723  341534041.0
```

[27]:
```
print('Mean life expectancy is', mean_life_expectancy)
print('75th Percentile life expectancy is', lower_boundary_high_life_expectancy)
```

```
Mean life expectancy is 72.75275
75th Percentile life expectancy is 77.937
```

[28]:
```
df_higher_gdp[df_higher_gdp['Life expectancy'] >␣
 ↪lower_boundary_high_life_expectancy]
```

[28]:

	Country	Year	Life expectancy	GDP per capita	\
3579	Australia	2022	82.766	51090.260	
3840	Austria	2022	81.296	55867.184	
4624	Bahrain	2022	80.992	51854.715	
5661	Belgium	2022	81.159	53287.152	
6273	Bermuda	2022	82.062	81165.650	
9564	Canada	2022	81.249	49296.380	
9879	Cayman Islands	2022	79.984	71353.890	
13156	Cyprus	2022	80.434	44996.316	
13417	Czechia	2022	79.165	41052.348	
14024	Denmark	2022	81.291	59935.120	
16671	Estonia	2022	78.167	37711.820	
18745	Finland	2022	81.243	49275.152	
19006	France	2022	82.475	45904.410	
20258	Germany	2022	80.580	53969.625	
23603	Hong Kong	2022	83.485	58478.883	
24125	Iceland	2022	81.588	55567.438	
25430	Ireland	2022	82.050	112445.420	

25767	Israel	2022	82.814	44393.300
26028	Italy	2022	82.052	44292.190
26542	Japan	2022	84.054	41837.910
27963	Kuwait	2022	78.788	49400.355
31506	Luxembourg	2022	82.201	117746.990
31592	Macao	2022	82.103	51840.140
33118	Malta	2022	82.250	48641.850
37552	Netherlands	2022	81.912	59249.168
38016	New Zealand	2022	81.006	45185.312
40106	Norway	2022	82.631	67296.160
43244	Portugal	2022	81.194	35767.723
43624	Qatar	2022	81.857	96557.810
47945	Singapore	2022	82.921	108036.110
48541	Slovenia	2022	80.793	41015.227
50005	South Korea	2022	82.727	45560.125
50493	Spain	2022	82.366	40223.010
51673	Sweden	2022	83.046	55359.344
51934	Switzerland	2022	83.200	72278.210
56201	United Arab Emirates	2022	80.487	74917.670
56462	United Kingdom	2022	81.074	47587.168
56723	United States	2022	77.979	64623.125

	Population
3579	26200987.0
3840	9064679.0
4624	1533459.0
5661	11641813.0
6273	64772.0
9564	38821260.0
9879	71609.0
13156	1331376.0
13417	10673216.0
14024	5902898.0
16671	1350092.0
18745	5569299.0
19006	66277412.0
20258	84086228.0
23603	7465914.0
24125	380368.0
25430	5110013.0
25767	9103144.0
26028	59619106.0
26542	124997586.0
27963	4589514.0
31506	653316.0
31592	704359.0
33118	528194.0

```
37552   17904422.0
38016    5131733.0
40106    5456795.0
43244   10417075.0
43624    2892465.0
47945    5649886.0
48541    2115230.0
50005   51782514.0
50493   47828386.0
51673   10487333.0
51934    8792180.0
56201   10242085.0
56462   68179315.0
56723  341534041.0
```

[29]:
```
df_higher_gdp[(df_higher_gdp['Life expectancy'] >= mean_life_expectancy) &␣
↪(df_higher_gdp['Life expectancy'] <= lower_boundary_high_life_expectancy)]
```

[29]:
```
                       Country  Year  Life expectancy  GDP per capita  \
2853                     Aruba  2022           76.226       41273.613
7998                    Brunei  2022           72.917       58669.902
23864                  Hungary  2022           76.212       35356.777
30715                Lithuania  2022           74.696       39955.246
40706                     Oman  2022           77.911       35336.895
42983                   Poland  2022           77.923       37706.605
46611             Saudi Arabia  2022           77.310       50188.297
48019  Sint Maarten (Dutch part)  2022           76.180       41942.918

       Population
2853     107792.0
7998     455374.0
23864   9684306.0
30715   2816922.0
40706   4730227.0
42983  38385734.0
46611  32175352.0
48019     42163.0
```

[30]:
```
df_higher_gdp[df_higher_gdp['Life expectancy'] < mean_life_expectancy]
```

[30]:
```
      Country  Year  Life expectancy  GDP per capita  Population
22553  Guyana  2022           69.888       35634.688    821636.0
```

[31]:
```
plt.figure(figsize=(18, 8))
plt.scatter(df_higher_gdp['GDP per capita'], df_higher_gdp['Life expectancy'],␣
↪color='blue', s = df_higher_gdp['Population']/200000, alpha=0.5)
```

```
for i, country in enumerate(df_higher_gdp['Country']):
    plt.text(df_higher_gdp['GDP per capita'].iloc[i], df_higher_gdp['Life␣
 ↪expectancy'].iloc[i], df_higher_gdp['Country'].iloc[i], fontsize=6,␣
 ↪ha='right')

plt.axhline(y=lower_boundary_high_life_expectancy, color='green',␣
 ↪linestyle='--', label='75th Percentile Life Expectancy')

plt.axhline(y=mean_life_expectancy, color='red', linestyle='--', label='Mean␣
 ↪Life Expectancy')

plt.xlabel('GDP per Capita (international-$ in 2017 prices)')
plt.ylabel('Life Expectancy (years)')
plt.title('GDP per Capita vs Life Expectancy (strong economies) for 2022')
plt.legend()
plt.show()
```